

Exercise 2 - Corona papers

The world is changing as we speak. The COVID-19 epidemic has been causing dramatic changes to the way we live our life. In this exercise we will focus on data related to the COVID-19 and will build unsupervised machine learning models.

The data that you will play with is a big corpus of academic papers, all dealing with COVID-19. The data is part of a worldwide research challenge, provided by Kaggle and is well described [here](#).

Your mission is divided into two parts:

1. Design a method that will be able to retrieve the K most similar instances to a given instance (instance in our case is an academic paper). In addition, your method is required to retrieve the actual distance found between the K instances and the given instance.
Your method must use **ONLY** the text data provided in the Kaggle competition. In addition, you **MUST** use, as the primal tool, ANY **compression** application or tool of your choice. Please note that you can plug (as a black-box) the compression application to a learning process, which in turn will act as a “meta” learning algorithm. The compression output will be used as the “distance” measure, to be furthered be used for returning the K most similar instances.
2. Use the method you designed and built in part 1 task id 573 in the [Kaggle competition](#) (in this Kaggle competition, few different tasks were introduced, all using the same dataset – we wish to focus on a single one in the exercise). Task id 573 is titled as [“What has been published about information sharing and inter-sectoral collaboration?”](#)
The exact way to use the method you built in part 1 for the purpose of this task is up to your choice. We encourage you to be creative and open minded.

Few important notes about the data maintained in the Kaggle challenge:

- Currently, there are ~40K full text academic papers in the website. Make sure you use only academic papers with full text (some of the academic papers that appear in the corpus contain only meta information about the paper).
- Data is being updated along time. You can download the current data available in the website and treat it as frozen fixed dataset (without updating it along time).
- Data is organized in a very structured way, but the structure is a bit complex. Take some time to understand the structure of the data. Then, make sure you have preprocessing functions that are able to pull out the required data for your main task.
- Feel free to use Kaggle forums (mainly the ‘Discussion’ and ‘Kernels’ sections in the website) for technical questions as well as for data science related questions. Indeed, you can be inspired by some of the solutions suggested in the website already.

Submission instructions

- The submission must contain your code (written in Python 3).
- Your code must contain a main file, as an entry point for others to use. As an alternative, you can add a notebook, including a general flow of your code (including as much explanations as possible). It must be documented in a reasonable way and contain simple instructions how it has to be used. If you use special Python packages (on top of what Anaconda provides) – please document it well.
- The submission must also contain a short document (up to 3 pages) describing the work you did and an analysis of your solution.
- The exercise has to be done individually.
- The exercise has to be submitted by 20.5.2020, 23:00.

We will take all of these elements in consideration when checking your submission. Our main focus will be on creativity and the usage of information theoretic considerations when designing your solution. Please submit even a partial solution in order to get a grade.